

Beacon — UX & PLG User Flows

v0.1 — UX/UI + Product-Led Growth Spec *Companion to: Beacon PRD v0.1*

1. Approach & Operating Principles

This document specifies the user flows that turn the Beacon product spec into a working PLG product. Two perspectives are covered:

- 1. **Customer-facing flows** — what Mike (owner) and his office manager experience from first touch through expansion and renewal.
- 2. **Internal admin flows** — what the Beacon team (operations, customer success, content review, billing) experiences running the platform.

Every flow is documented with: goal, entry trigger, step-by-step screen flow, mobile-first wireframe of the critical screen, PLG levers, failure modes and recovery, and the event instrumentation needed to measure it.

PLG framework guiding every customer flow

Principle	What it means in Beacon
Time-to-Value < 2 minutes	Owner sees their AI Visibility Score before they finish their coffee
Concrete value before payment	Free audit produces a real, personalized score; first fix free
Self-serve activation, assisted only when stuck	No mandatory onboarding call; CSM intervenes only on explicit signal
Progressive disclosure	First week: one button (approve). Month three: full power-user capability available but not required
Earned automation	Approval mode default; autopilot offered only after trust signals fire
Compound habit, not dashboard fatigue	Weekly SMS digest replaces the "log in and check" model
Reversibility everywhere	Every published change has a one-tap revert; cancellation never destroys shipped assets

Principle	What it means in Beacon
Distribution baked into the product	Score certificate, peer-comparison, agency white-label all live inside the product

UX principles guiding every screen

- 1. **Two screens.** Home + Action Queue. Everything else is sub-views.
- 2. **One number on Home.** AI Visibility Score, big, with a 30-day sparkline.
- 3. **No jargon, ever.** "How often AI mentions you," not "share of voice."
- 4. **Big tap targets.** Mike approves from his phone with grease on his hands.
- 5. **Single-color status.** Green (working), yellow (needs attention), red (broken).
- 6. **Mobile-first PWA.** Native-app feel, web cost.
- 7. **Undo > Confirm.** A revert button beats five confirmation dialogs.

2. Personas & State Definitions

Customer-side personas


Persona	Role	Where they live in the product
Mike (Owner)	Buyer, decision-maker, license holder	Home screen, weekly digest, billing, approval queue
Sarah (Office Manager)	Day-to-day operator, often the "champion"	Approval queue, review reply, weekly report
Field Tech (rare)	Receives review-request links via SMS	Mobile job-complete trigger only

Internal-side personas

Persona	Role	Where they live
Operations Lead	Customer health, churn, expansion	Customer Health Dashboard, churn alerts
Content Reviewer	First 30 days of every customer; ongoing low-confidence items	Content Review Queue

Persona	Role	Where they live
CSM (post-MVP)	Pro+ customer relationships	Account view, QBR notes, expansion playbook
Support	Tier-1 tickets	Inbox + macros
Agency Admin (V2)	Marketing agency reselling Beacon	Multi-tenant admin

Customer lifecycle states



Each state has a **definition**, an **entry event**, and an **exit/promotion event**:

State	Definition	Entry event	Promotion event
Visitor	Landed on beacon.app	<code>pageview_landing</code>	<code>audit_started</code>
Audit Completed	Saw their AI Visibility Score	<code>audit_completed</code>	<code>signup_started</code>
Trialing	Signed up, in 14-day Pro free trial	<code>trial_started</code>	<code>first_fix_shipped</code>
Activated	Approved + shipped at least one fix	<code>first_fix_shipped</code>	<code>payment_method_added</code>
Paid	Billing live, subscription active	<code>subscription_active</code>	<code>weekly_digest_opened_4_weeks</code>
Habituated	Engaged with weekly digest 4 weeks running	<code>digest_streak_4</code>	<code>autopilot_enabled</code>
Autopilot	Owner toggled on full autonomy	<code>autopilot_enabled</code>	<code>tier_upgraded</code> OR <code>referral_sent</code>
Expanded	Upgraded tier or invited team member	<code>tier_upgraded</code> OR <code>seat_added</code>	—
At-Risk	Score plateau, missed digests, low approval	<code>health_score_yellow</code>	back to Habituated OR Churned
Churned	Subscription canceled	<code>subscription_canceled</code>	<code>winback_succeeded</code>

Every flow below is mapped to one or more of these states.

PART I — CUSTOMER-FACING FLOWS

Flow 1 — The Free Audit (PLG Hook)

State transition: Visitor → Audit Completed **Goal:** Convert anonymous traffic to a personalized, share-worthy result that creates buying intent. **Entry trigger:** Direct nav to `beacon.app`, paid ad, IEC chapter referral, or "audit your AI visibility" CTA from anywhere.

Why this is the most important flow in the product

Without signup, without payment, without a dashboard — the audit produces a real, personalized AI Visibility Score that **the contractor can immediately verify by typing the same prompts into ChatGPT himself**. That verifiability is the entire conversion engine. Everything else flows from this two-minute experience.

Step-by-step

Step 1 – Landing

beacon.app

See how AI search engines describe
your shop right now. Free, 90 seconds.

Your Google Business Profile or shop
name: Lone Star Electric, Plano TX

[See My AI Visibility Score →]

No signup required.

No credit card. No spam.

Step 2 – Auto-resolution (background, ~3 sec)

- Pull GBP data via Places API
- Pull website + sitemap
- Detect services from GBP categories + website
- Detect service area from GBP service-area zips
- Identify top-3 competitors via Maps proximity + category match

Step 3 – Confirmation (10 sec)

Is this you?

Lone Star Electric

3401 W 15th St, Plano TX 75075

Electrical contractor • 4.7★ (87 reviews)

Top services we'll check for:

- ✓ Panel upgrade
- ✓ EV charger install
- ✓ Generator install
- ✓ Electrical repair
- ✓ Whole-home rewiring

[+ Add] [Edit]

We'll compare you against:

- Plano Electric Pros
- Reliable Electric of North Dallas
- Bright Spark Electrical

[Edit competitors]

[Run My Audit →]

Step 4 – Live progress (~90 sec)

Asking AI engines about you...

- ✓ ChatGPT – "best electrician in Plano"
- ✓ Gemini – "panel upgrade Plano cost"
- ⌚ Perplexity – "EV charger install Plano"
- ⌚ ChatGPT – "emergency electrician Plano"
- ⌚ ...

 18 of 25 prompts

Did you know? 45% of homeowners now use ChatGPT to find local services. AI engines only recommend ~1.2% of local businesses.

Step 5 – RESULTS (the aha)

Your AI Visibility Score

12 out of 100

Out of 25 questions homeowners ask AI about electrical services in Plano:

- ✓ You appeared in 2
- Partially mentioned 1
- × You were absent in 22

Plano Electric Pros appeared in 8
Bright Spark Electrical appeared in 6

Top 3 reasons you're losing:

1. Plano Electric has 12 FAQ pages covering panel, EV, generator. You have none.
2. Bright Spark has city-specific service pages for Plano, Frisco, Allen. You only have one generic page.
3. Your GBP has no Q&A section. AI engines pull heavily from GBP Q&A.

- ▶ See all 25 prompts (free signup)
 - ▶ Ship your first fix this week, free

[Sign up free →] [Email this to me]

Key UX decisions (and why)

- **Single input field on landing.** Lifted from Linear/Vercel. Anything more is friction. We can derive everything else from a Google Maps lookup.
- **No signup before results.** The result is the demo. If we gate it behind email, we lose 60% of audit completers. The signup happens *after* they see the gap.

- **Live progress with rotating "did you know" copy.** Every second of waiting needs to build value. Educational copy doubles as soft pitch.
- **Specific competitor names rendered.** This is provocative — Mike feels the gap viscerally when he sees Lone Star vs. Bright Spark in numbers. Avoiding this is being too polite to convert.
- **Reasons shown in plain language.** "Plano Electric has 12 FAQ pages, you have none" — not "FAQ schema density delta of 0.84."
- **Two CTAs.** "Sign up free" is primary; "Email this to me" is the soft escape that still captures lead and creates a follow-up touchpoint.

PLG levers

Lever	Mechanism
Zero-friction first-touch	No signup required to see result
Concrete personalized output	Owner can verify any prompt in ChatGPT himself
Emotional provocation	Specific competitor names + numbers
Soft commitment escalation	"First fix free" makes signup feel like collecting earned value
Latent virality	"Email this to me" → owner forwards to office manager / business partner / accountant
Score certificate (V1.5)	Shareable image: "My AI Visibility Score is 87. What's yours?" — viral on LinkedIn

Failure modes & recovery

- **Business not found in GBP.** Show a manual entry form with website URL field. Take 30 sec longer; works.
- **No competitors auto-detected.** Default to category-level competitors ("we'll compare you against the top 3 electrical contractors within 10 miles").
- **All 3 LLM APIs slow or down.** Use 24h-cached results; show "based on this morning's data" footnote. Never break the flow.
- **Owner gets a high score (>50).** Reframe: "You're ahead — let's make you uncatchable." Same conversion, different angle.

- **Owner gets a very high score (>80).** Honest framing: "You're in the top 1% already. We'd add maintenance and competitor watch — but you may not need us. Want to see what we'd do?"

Events to instrument

```
audit_landed           { source, utm_*, device }
audit_business_entered { business_name, gbp_resolved: bool }
audit_confirmed        { services_count, competitors_count }
audit_running          { prompt_count, engines: [...] }
audit_completed        { score, time_to_complete_seconds }
audit_results_viewed   { time_on_page_seconds }
audit_signup_clicked    { from_section: 'cta' | 'banner' }
audit_email_captured    { variant: 'send_to_self' }
audit_shared           { channel: 'linkedin' | 'sms' | 'email' }
audit_competitor_swapped { from_auto, to_manual }
```

Success metrics

- Audit completion rate: target **75%** (visitor → result)
- Audit → signup rate: target **25%**
- Audit → email-capture rate: target **40%** (combined funnel for nurture)
- Time-to-result: target **<120 seconds**
- Audit-shared rate: target **5%** (viral coefficient seed)

Flow 2 — Signup → Onboarding → First Aha

State transition: Audit Completed → Trialing → Activated **Goal:** Get to "first fix shipped" within 48 hours of signup, ideally <24h. **Entry trigger:** "Sign up free" CTA from audit results.

Step-by-step

Step 1 – Account creation (15 sec)

```
| Create your Beacon account |
|                             |
| Continue with Google [G]   [or use email] |
|                             |
| We'll save the audit you just ran. |
| No credit card needed for the next 14 days. |
```

Step 2 – One-screen confirmation (audit data pre-filled, 30 sec)

Quick check – these are pre-filled from your audit. Edit anything that's wrong.	
Business: Lone Star Electric	[🔗]
Service area: Plano + 6 nearby zips	[🔗]
Top services: 5 services	[🔗]
Competitors tracked: 3	[🔗]
Brand voice samples we'll learn from:	
✓ 87 Google reviews	
✓ Your About Us page	
✓ 12 GBP posts from past 6 months	
[Add more (LinkedIn, ads, voicemail)]	
Unlock your full radar (50 prompts). Then we'll ship your first fix free.	
[Run my full radar →]	

Step 3 – Full radar runs (~90 sec, background)

- All 50 prompts × 3 engines
- Brand voice profile built from collected samples
- Initial gap inventory generated

Step 4 – Welcome to Home screen (the aha)

- Score card animates in
- Three "this week" cards populate
- Top notification: "Your first fix is ready to review"

Step 5 – First-Fix Preview (within 5 min of signup)

Your first fix – free, ready to ship
WHAT WE'LL ADD
6 frequently-asked-question pairs to your panel-upgrade page.
WHY THIS
ChatGPT cited Plano Electric Pros for

'panel upgrade Plano cost' – they have a detailed FAQ. You don't have one.

PREVIEW (in your voice)

Q: How much does a 200-amp panel upgrade cost in Plano, TX?

A: In Plano and surrounding North Dallas suburbs, a 200-amp panel upgrade typically runs \$2,800-\$4,500, depending on whether your home...

...– Lone Star Electric, licensed Master Electrician, serving Plano since 2009.

+ 5 more FAQ pairs [See all 6]

WHERE WE'LL PUBLISH

✓ lonestarelectric.com/panel-upgrade

✓ Your Google Business Profile Q&A

[Change]

Want us to first install our 1-line script tag on your site? 5 minutes.

[Skip – manual paste] [Install now →]

Step 6 – Schema script install (or skip)

- Walkthrough screencast for WordPress / Webflow / Wix
- For "other": copy-paste script tag, verify on site
- Verification check confirms it's live

Step 7 – Ship the first fix

- Owner taps Approve
- "Live in 5 minutes" loading screen
- Confirmation: live URL + "We'll re-ask ChatGPT in 7 days"

Step 8 – End of onboarding state

- Lands on Home screen
- Banner: " First fix live. Your 14-day trial just started."

- "What's next" tooltip points to Action Queue: "We'll keep finding gaps. Approve them as they come."

Key UX decisions

- **No "set up your account" multi-page wizard.** Everything from the audit pre-fills. The signup is a single confirmation screen.
- **Brand voice samples are auto-collected from public sources.** Owner can add private sources (LinkedIn, ads, voicemails) but nothing is required.
- **First-fix preview shows the actual content, not a placeholder.** This is the credibility moment. If the FAQ sounds generic, we've lost the customer.
- **Schema script install is offered, not required.** Skipping is fine — the first fix can publish to GBP only. Site fixes can wait until the script is installed.
- **End-state is the Home screen, not "you're all set!" success page.** Drop the user into the product, not a marketing wrapper.

PLG levers

Lever	Mechanism
Pre-filled wizard	Reduces signup-to-aha to <90 seconds
Free first-fix	Inverted commitment — they receive value before they pay
14-day trial without credit card	Removes payment friction
Brand voice from public sources	Demonstrates product-aware sophistication; no extra work for owner
Live URL after first publish	Owner can verify; shareable proof
Banner on Home: "Trial just started"	Transparency about state, no surprise charges

Failure modes & recovery

- **Owner abandons during radar load (90 sec).** Send email + SMS in 5 min: "Your radar is ready — see your score." With deep link.
- **First-fix preview content sounds wrong.** Edit-mode is rich text, not form-based. Allow inline edits with one-click "regenerate this paragraph."

- **Owner refuses script install AND no GBP connected.** Allow first fix to go to a Beacon-hosted preview page (`([shop].beacon.site/panel-upgrade-plano)`) so they can see the asset and we can demonstrate value, even before they commit to publishing on their own site.
- **Owner doesn't approve first fix within 24h.** SMS reminder: "Your first fix is still ready. 30 seconds to review and ship it."
- **Owner doesn't approve first fix within 72h.** CSM gets alert (Pro+ tier).

Events to instrument

```

trial_started           { source: 'audit' | 'direct' }
onboarding_step_completed { step: 1..7 }
brand_voice_samples_added { sources: [...], total_samples }
schema_script_installed  { method: 'wordpress' | 'webflow' | 'wix' |
'manual', verified: bool }
first_fix_previewed      { time_since_signup_seconds }
first_fix_edited         { paragraphs_changed }
first_fix_approved       { time_since_signup_seconds }
first_fix_shipped        { fix_type, channels: ['site', 'gbp'] }
onboarding_abandoned    { last_completed_step, idle_seconds }

```

Success metrics

- Signup → first-fix-shipped rate: target **70% within 48h**
- Time to first fix (median): target **<3 hours**
- Schema-script install rate at onboarding: target **45%** (acceptable to defer; not required for activation)
- Onboarding abandonment: target **<15%**

Flow 3 — Daily Habit Loop (Passive)

State: Activated → Habituated **Goal:** Continuous Score lift + content shipped without owner attention.

This is not a flow the owner *does* — it's a flow that *happens*. The owner's only interaction is the weekly digest (Flow 4) and approval queue (Flow 5).

What runs daily

```

2:00 AM    Visibility Radar runs (50 prompts × 3 engines for every active customer)

```

```
2:30 AM    Citation parser: identifies who got cited for what
3:00 AM    Gap diagnosis: for every gap, fetch competitor's source and run LLM diff
4:00 AM    Fix generation: queued items pre-generated and brand-voice-checked
6:00 AM    Score recomputation per customer
6:30 AM    Day's queue ready for owner review (or autopilot publish)
8:00 AM    Catchup: anything that didn't complete by 6 AM
```

The owner sees zero of this. The product is doing work overnight; he wakes up and the queue is fresh.

Edge cases handled silently

- **LLM API rate limits hit.** Queue spills to next day; alert internal ops only if persistent.
- **Customer's website is down.** Skip site-publish actions; queue them for retry. Notify owner only if down for 48+ hours.
- **GBP API quota hit.** Spread posts across the week; never miss the cadence.
- **A scheduled publish fails.** Auto-retry 3 times over 48h; surface to internal queue if all fail.

Flow 4 — Weekly Digest (The Habit)

State: Habituated maintenance **Goal:** Replace the "log in and check" model with a passive 60-second weekly read. **Entry trigger:** Sunday 6 AM (configurable) — SMS + email send.

The SMS

```
[SMS — Sunday 6:04 AM]
Beacon weekly: Score 31 → 38 (+7).
4 fixes shipped. You appeared in 3 new
ChatGPT answers — incl. "best EV charger
install in Plano." 5 fixes in queue.
Tap to review: bcn.ai/w/aB3
```

The email (richer)

Subject: Lone Star Electric — AI Visibility +7 this week

```
| WEEK OF NOVEMBER 24, 2026 |
```

```
| Your AI Visibility Score |
```

31 → 38 (+7 this week)

 (last 12 weeks)

WHAT WE SHIPPED

1. Added 6 FAQs to your panel-upgrade page
2. Published a Frisco-specific service page (long-tail target)
3. Replied to 4 reviews in your voice
4. Refreshed your GBP Q&A section with 8 entries

WHERE YOU NOW APPEAR

- ChatGPT now mentions you for "best EV charger install in Plano"
- Perplexity cites your panel-upgrade page directly in answers
- Gemini mentions you in 2 new prompts

WHAT'S QUEUED FOR NEXT WEEK

- 3 service pages (Allen, McKinney, Wylie)
- 4 review replies
- 1 weekly GBP post

[Approve all (1 tap)]

[Review individually]

Score breakdown:

Citation rate 12/50 (+3)

Sentiment strong

Share of voice 18% (+4%)

Citation diversity 11 sources (+2)

Compared to your 3 tracked competitors, you're now ranked: #2 (was #4).

Key UX decisions

- **SMS is the trigger; email is the depth.** Most owners open the SMS first; some click into

email. Both lead to the same web page.

- **One number in the subject line.** "+7 this week" — clear value signal.
- **"Approve all" is a single tap.** This is the engine of autopilot graduation.
- **Competitor ranking shown.** Triggers competitive emotion → drives engagement.
- **No "ROI in dollars" framing in the digest itself.** Reserve that for monthly business reviews — weekly digests should celebrate the work, not over-promise outcomes.

PLG levers

Lever	Mechanism
SMS over email	95%+ open rate vs. 25% for email
Sunday morning send	Off-hours read, calm context
" +N this week" hook	Concrete weekly progress
One-tap approve all	Major step toward autopilot
Competitive ranking	Emotional engagement → retention

Failure modes & recovery

- **Score went down this week.** Honest framing: "Score 38 → 35 (-3). Two competitors shipped new content. We're working on it. Here's what we'll do this week." Trust > spin.
- **No fixes shipped this week.** This shouldn't happen, but if it does: "Catching up on diagnostics. Full report next week." Internal alert fires.
- **Owner stopped reading digests.** After 3 missed weeks: "We notice you haven't opened your digest. Want to switch to monthly? Or pause autopilot?" — proactive, not nagging.

Events to instrument

digest_sent	{ channel: 'sms' 'email', score, score_delta }
digest_opened	{ channel, time_to_open_minutes }
digest_link_clicked	{ destination: 'web' 'approve' }
digest_approve_all_tapped	{ items_approved }
digest_unsubscribed	{ reason }
digest_streak_4_weeks	(state transition event)

Success metrics

- Digest open rate: target >**80%** (SMS) / >40% (email)
- Approve-all tap rate: target >**50%** of opens
- 4-week digest streak: target >**60%** of paid customers (Habituated state)

Flow 5 — Action Queue & Approval Flow

State: Activated maintenance / Habituated **Goal:** Mike approves changes from his phone with grease on his hands. <30 sec per approval.

Mobile screen — Action Queue

← Action Queue	Sort ▼
5 pending	1 high-priority
Reply to negative review Susan T., 2 stars, posted 14h ago "Tech showed up late, no call..." Our draft reply → [Preview] [Approve] [Edit]	
Add 6 FAQs to panel-upgrade page Why: ChatGPT cited Plano Electric for "panel cost Plano" – not us [Preview] [Approve] [Edit]	
Publish service page: Frisco TX Long-tail: "EV charger install Frisco" – no current presence [Preview] [Approve] [Edit]	

GBP weekly post: AFCI explainer

[Preview] [Approve] [Edit]

Refresh About Us schema markup

[Preview] [Approve] [Edit]

[Approve all 5] [Skip all]

Want to skip approvals?

Turn on Autopilot → (you've approved 96%
of past items – you're ready)

Tapping "Preview" on an item

← Preview

Approve

Reply to Susan T. (2 stars)

—— HER REVIEW ——

Susan T. 14h ago

"Tech showed up late, no call ahead.
Work was fine but really frustrating."

—— OUR DRAFT REPLY ——

Susan, I'm Mike at Lone Star and I
want to apologize for the late arrival
– that's on us. Please call me direct
at 972-xxx-xxxx so I can make this
right. Thanks for letting us know.

► Why this approach: Acknowledging late
arrival without excuses, owner's name,
direct phone number, ownership tone.

[Edit text] [Regenerate]

[APPROVE & POST]

Once approved, this posts publicly to Google. You can edit it after if needed.

Key UX decisions

- **Cards, not rows.** Each item is a self-contained card with all context inline. No need to drill into a detail screen for a quick approve.
- **Color-coded urgency.** = needs attention (negative review, urgent), = standard.
- **Why this? rendered inline.** Owner doesn't trust changes he doesn't understand. Two-line "why" is required on every card.
- **Preview vs. Approve are different actions.** Approve commits; Preview shows.
- **"Approve all" is the biggest button at the bottom.** PLG-critical for autopilot graduation.
- **Autopilot graduation prompt at the bottom.** Earned, not pushed: "you've approved 96% of past items."

PLG levers

Lever	Mechanism
Mobile-first design	Approval from the truck cab
Single-tap approve	Reduced friction → higher engagement → habit formation
Approve-all button	Direct path to autopilot
Autopilot graduation prompt	Behavioral trigger when approval rate $\geq$ 90%
Preview before commit	Trust mechanism; reduces revert rate

Failure modes & recovery

- **Owner approves something he didn't mean to.** One-tap revert from confirmation screen, also from Action Log (Settings → History → Revert).
- **Owner edits AI draft and ships a typo.** Revert + re-edit; we don't punish him for it.
- **Owner skips items repeatedly.** Internal alert: voice profile may be off. CSM intervention.

- **Negative-review reply needs human review (legal-sensitive language).** Auto-routes to internal Content Reviewer queue (admin flow #12) before showing to owner.

Events to instrument

```
queue_viewed           { items_count, high_priority_count }
item_previewed         { item_id, item_type }
item_approved          { item_id, time_in_preview_seconds }
item_edited            { item_id, edit_chars }
item_regenerated        { item_id }
item_skipped           { item_id, reason? }
item_reverted          { item_id, time_since_publish_seconds }
approve_all_tapped     { items_count }
autopilot_prompt_shown { approval_rate, items_history }
autopilot_enabled      { from_prompt: bool }
```

Success metrics

- Approval rate: target >**85%** (skipped + reverted < 15%)
- Time-per-approval: target <**30 sec** median
- Edit rate: target <**10%** (>10% means voice profile is off)
- Revert rate: target <**5%** (>5% means quality is failing)

Flow 6 — Autopilot Graduation

State transition: Habituated → Autopilot **Goal:** Get the owner out of the queue. Frees mental load. Drives retention. **Entry trigger:** Behavioral signal (approval rate $\geq$ 90%, 30+ days, no reverts in 14 days).

The graduation moment

```
| You're ready for Autopilot. |
|                               |
| In your first 35 days you've approved |
| 47 of 51 fixes. Edits: 3. Reverts: 0. |
|                               |
| You can let Beacon ship straight to live. |
|                               |
| WHAT CHANGES |
|   ✓ No more approval queue (Beacon ships |
```

```
|   directly) |
|   ✓ Weekly digest still arrives Sundays |
|   ✓ One-tap revert on anything |
|   ✓ Pause autopilot anytime |
| |
| WHAT STAYS THE SAME |
|   ✓ All shipped content stays on your site |
|   ✓ Negative reviews still need approval |
|   ✓ Major site changes still need approval |
| |
| [ Turn on Autopilot ] |
| [ Not yet, keep approving ] |
```

Key UX decisions

- **Earned, not pushed.** Don't show this prompt until the behavioral threshold is met. Showing it too early kills trust.
- **Negative reviews always require approval.** This is non-negotiable. Even on autopilot, certain content types route to manual review.
- **Reversibility in plain language.** "Pause autopilot anytime" + "one-tap revert" — fear-killers.

PLG levers

Lever	Mechanism
Behavioral trigger	Only shown when ready → high acceptance rate
Reversibility framing	"Pause anytime" reduces commitment fear
Statistical proof	"You approved 47 of 51" → owner sees his own readiness
Cognitive offload	After enabling: owner stops thinking about Beacon → strongest retention signal

Events to instrument

```
autopilot_eligible      (behavioral state event)
autopilot_prompt_shown  { trigger_context }
autopilot_enabled       { time_since_signup_days }
```

autopilot_paused	{ reason }
autopilot_re_enabled	{ time_paused_days }

Success metrics

- % of paid customers on autopilot at month 3: target **40%**
 - % of paid customers on autopilot at month 6: target **70%**
 - Autopilot pause rate: target **<10%** (rare, indicates quality problem)
-

Flow 7 — Score Plateau / Recovery

State transition: Habituated → At-Risk → Habituated **Goal:** Detect when the product has stopped delivering Score lift and intervene before churn. **Entry trigger:** Score flat or declining for 3 consecutive weeks.

What the owner sees (gentle, honest)

[SMS]
Heads up — your AI Visibility Score has been flat at 42 for 3 weeks. We're going to run a deeper diagnostic and try a different angle. Update next week. Reply CALL if you want to talk to a human.

What runs in the background

- Internal alert fires; assigned to CSM (Pro+) or to ops bench (Lite).
- Deep diagnostic: re-fetch all competitor cited content, look for new patterns we missed, re-evaluate prompt set.
- Different fix types queued (e.g., shift from FAQ generation to authority-building external citations).
- Owner gets next weekly digest with "we tried X, here's what's working / not."

Key UX decisions

- **Honest framing, not corporate spin.** Trades audience smells corporate-speak instantly. "Heads up — we're working on it" beats "We've identified opportunities."
- **Reply-to-talk option.** SMS reply with "CALL" triggers a CSM callback within 4h. No support form, no portal.

- **Next digest closes the loop.** Don't leave the conversation hanging.

Failure modes & recovery

- **3 more weeks of flat Score after intervention.** Escalate: free strategy session offered, or proactive 30%-off-for-3-months retention offer.
 - **Owner replies CALL.** CSM hits 4-hour SLA. CSM has access to full diagnostic data before the call.
 - **Score plateau is genuine ceiling (already at top).** Reframe: shift to maintenance mode + competitor watch. Score isn't going up because there's nowhere to go.
-

Flow 8 — Team Member Invite

State: Habituated / Autopilot **Goal:** Mike brings his Office Manager into the product.

Step-by-step

Settings → Team → [+ Invite member]

Invite a teammate

Email: sarah@lonestar.com

Role:

- ☐ Owner (full access incl. billing)
- ☒ Manager (queue, reports – no billing)
- ☐ Viewer (read-only)

[Send invite]

Manager's experience

- Receives email with magic-link sign-in.
- Lands on shared workspace home.
- Sees same Action Queue Mike sees.
- Cannot access Settings → Billing or Settings → Integrations (changes require Owner).
- Can approve, edit, regenerate items.

- Can configure her own digest preferences (digest can go to multiple addresses).

Key UX decisions

- **Magic-link signup for invitees.** No password creation friction.
- **Role-based defaults.** Manager doesn't see billing. Reduces accidental confusion.
- **Shared workspace, not separate accounts.** Owner and Manager work on the same data.

PLG levers

- **Multi-seat = stickier account.** When Sarah uses the product daily, Mike doesn't churn it.
- **Sarah may become the in-product power user.** Owners often disengage; champions retain accounts.

Flow 9 — Tier Upgrade (Lite → Pro)

State transition: Paid (Lite) → Paid (Pro) = Expanded **Goal:** Upgrade triggered by hitting a value ceiling or unlocking a gated feature.

In-context upgrade triggers

Trigger	Mechanism
Lite user hits 25-prompt cap	Banner: "You're tracking 25 of 50 possible prompts. Pro unlocks all of them."
Lite user views autopilot toggle	"Autopilot is a Pro feature. Try it free for 7 days."
Lite user wants a Service Page generated	"Service Page generation is on Pro. See pricing →"
Score plateau on Lite	"Lite is monitoring-only. Pro ships fixes automatically — that's where most Score lift comes from."

Upgrade modal

Upgrade to Pro		
	LITE	PRO
Price	\$199/mo	\$499/mo

Prompts tracked	25	50
Auto-publish	x	✓
Schema injection	x	✓
Service pages	x	✓
Brand voice	basic	deep
FSM integration	x	✓
Slack digest	x	✓

In your first 60 days on Lite, your Score went 12 → 24. Pro customers in similar shops average 12 → 47 in 60 days.

[Upgrade – try Pro free for 7 days]
 [Stay on Lite]

Key UX decisions

- **Comparison is concrete, not feature-fetishy.** "Pro customers in similar shops average 12 → 47 in 60 days" beats a feature checklist.
- **7-day Pro trial when upgrading from Lite.** Reduces decision friction.
- **Downgrade always available.** No "Pro only" lock-in.

PLG levers

- **In-context upgrades.** Triggered exactly when the value gap is felt.
- **Cohort-based proof.** "Customers like you" comparison is more persuasive than abstract feature lists.
- **Trial-to-trial.** Upgrade is itself a trial — reduces conversion fear.

Events to instrument

```

upgrade_prompt_shown      { trigger, current_tier }
upgrade_modal_viewed     { trigger }
upgrade_clicked           { from_tier, to_tier }
upgrade_completed        { time_since_prompt_minutes }
upgrade_declined          { reason? }
upgrade_trial_started     { from_tier, to_tier }

```

Flow 10 — Referral / Peer Loop

State: Habituated / Autopilot / Expanded **Goal:** Make peer recommendation in IEC chapters and supply houses easy and rewarded.

Referral mechanism

Settings → Refer a friend

Refer a peer, get \$100

Your unique link:

beacon.app/r/lonestar

[Copy] [Share via SMS] [Email]

Or send their score directly:

Friend's business: _____

[Run their audit & email it]

——— YOUR REFERRALS ———

- Bright Spark Electrical – joined Sep 14
You earned \$100 credit
- Plano Electric Pros – running audit

Score certificate (V1.5)

After 90 days on Beacon, owners get a shareable image:

AI VISIBILITY SCORE

73

▲ +61 in 90 days

LONE STAR ELECTRIC

Plano TX • 87 reviews • Master licensed

#1 ranked AI-cited electrician in Plano

Powered by Beacon

[Download] [Share LinkedIn] [Share GBP]

Key UX decisions

- **Two paths: link sharing OR direct audit-email.** The audit-email path is more powerful — Mike can run a peer's audit and forward the result, planting the seed personally.
- **Concrete dollar incentive.** \$100 credit (per side, both Mike and the referred peer get \$100).
- **Score certificate is shareable.** Bragging rights → viral coefficient.

PLG levers

Lever	Mechanism
Two-sided incentive	Both parties motivated
Audit-as-gift	Mike sends a peer their personalized audit — relationship-grade lead gen
Score certificate	LinkedIn/GBP shareable artifact → top-of-funnel for Beacon
IEC chapter / supply house context	Trades community is tight-knit, peer recs convert

Events to instrument

```
referral_link_copied
referral_link_shared      { channel }
referral_audit_sent       { peer_business }
referral_signup_attributed { referrer_id }
referral_conversion        { referrer_id, peer_id, credit_amount }
score_certificate_downloaded
score_certificate_shared    { channel }
```

Success metrics

- Referral coefficient: target **0.15** (each customer brings 0.15 paid customers)
 - Score certificate share rate: target **20%** of eligible customers (90+ days, Score ≥ 50)
-

Flow 11 — Cancellation & Win-Back

State transition: Paid $\rightarrow$ Churned (or back to Paid) **Goal:** Honest cancellation flow that preserves dignity, captures real reasons, and surfaces last-chance offers without dark-pattern friction.

Cancellation flow

Settings $\rightarrow$ Billing $\rightarrow$ Cancel subscription

Sorry to see you go.

Quick – what's the reason?

(This actually shapes our offer below.)

- ☐ Too expensive
- ☐ Not seeing enough value
- ☐ Selling/closing the business
- ☐ Switching to a different tool
- ☐ Score is great, don't need ongoing
- ☐ Other

[Continue]

Per-reason offer screen

Reason	Offer
Too expensive	"We get it. Want 50% off for 3 months while we work this out together?"
Not seeing enough value	"Score may be plateauing. Free 30-min strategy session with our lead?"
Selling/closing	"Best of luck. We'll keep your shipped content live for 90 days at no cost."
Switching	"Curious — which tool? Let us beat or match."
Score is great	"You're right — let's drop you to monitoring-only at \$79/mo. No more changes; we just watch."
Other	Free-text + 20% off offer

Final confirmation

Cancellation confirmed.

Here's what happens:

- ✓ Everything we built stays live on your site forever.
- ✓ Your last digest goes out Sunday.
- ✓ Schema script keeps working until you remove it.
- × We stop running radar and shipping new fixes after Dec 14.

Change your mind in the next 30 days and we'll restore everything as if you never left.

[Done] [Actually, never mind]

Win-back sequence (post-cancel)

- Day 7: Final monthly report shows what was working
- Day 14: New competitor citation alert ("Bright Spark just got cited 3 times this week")

- Day 30: "Score has dropped from 73 to 61 since you canceled. Want to come back?"
- Day 60: Last touch — case study from a similar peer
- Day 90: Stop the sequence

Key UX decisions

- **Reason captured first, offer shaped to it.** Dark patterns this isn't — but smart segmentation is fair.
- **Honest framing about what stays/goes.** "Everything stays live forever" is a real moat *and* a real reassurance.
- **30-day "as if you never left" window.** Reduces fear of leaving wrong.
- **Win-back uses real competitive data, not nagging.** "Your competitor pulled ahead" beats "we miss you."

PLG levers

Lever	Mechanism
Reason-shaped offers	Higher save rate without manipulation
Compounded asset preservation	Reduces psychological cost of leaving (paradoxically reduces actual leaving)
Real-data win-back	"Score dropped" is provable, not emotional

Events to instrument

```

cancellation_started
cancellation_reason_selected { reason }
retention_offer_shown       { offer_type }
retention_offer_accepted    { offer_type }
cancellation_completed      { reason, mrr_lost }
winback_email_sent          { day, content_type }
winback_clicked             { email_id }
winback_succeeded           { time_since_cancel_days }

```

Success metrics

- Cancel-to-save rate: target **25%** (1 in 4 reach the offer screen and stay)
- Win-back rate at day 90: target **8%** of churned customers

- Reason-coverage rate: target >95% (always capture *why*)

PART II — INTERNAL ADMIN FLOWS

These flows describe what the Beacon team does inside the platform.

Flow 12 — Content Review Queue (Quality Gate)

Persona: Content Reviewer **Goal:** Every AI-generated piece for new customers (first 30 days) and any low-confidence piece thereafter passes through human review before publishing.

Content reviewer dashboard

Content Review Queue247 pending

Filter: ▾ All customers ▾ All types ▾ Confidence < 0.7

Lone Star Electric (Day 3)

TYPE: Service page (Frisco)CONFIDENCE: 0.62

FLAGS: low brand-voice match; 1 phrase resembles training

[View]

[Approve]

[Edit & approve]

[Reject]

Bright Spark (Day 17)

TYPE: Negative review replyCONFIDENCE: 0.78

FLAGS: legal-sensitive language detected ("liable")

[View]

[Approve]

[Edit & approve]

[Reject]

Reliable Electric (Day 22)

TYPE: FAQ pairCONFIDENCE: 0.91

FLAGS: none

[View]

[Approve]

[Edit & approve]

[Reject]

Bulk actions: [Approve all confidence $\geq$ 0.85]

What "Reject" does

- Item never reaches the customer's queue.
- Generation is retried with adjusted parameters (different few-shot examples, voice profile re-tuning).
- After 2 rejections on same item type for same customer, internal alert: voice profile may need retraining.

Pattern detection

- If a customer has >5 rejected pieces in 30 days, the system flags their brand voice profile for re-evaluation.
- If a content type (e.g., FAQ pairs) has >15% rejection rate across all customers, the system flags the prompt template for re-engineering.
- If a single reviewer rejects >3x the rate of peers, internal QA reviews their decisions.

Key UX decisions

- **Confidence score visible.** Reviewer triages high-confidence items fast, focuses brain on edge cases.
- **Flags called out explicitly.** "Legal-sensitive language detected" pre-loads the reviewer's attention.
- **Bulk approve for high-confidence.** Reviewer time is finite; system should auto-approve trustworthy items.
- **Edit-and-approve preserves the customer experience.** The customer never sees that their item was reviewed by a human.

Events to instrument

review_queued	{ customer_id, content_type, confidence, flags }
review_approved	{ reviewer_id, time_in_review_seconds }
review_edited_approved	{ reviewer_id, edit_chars }
review_rejected	{ reviewer_id, reason }
review_pattern_alert	{ pattern_type, customer_id }
voice_profile_retrain_flagged	{ customer_id, trigger }

Success metrics

- Reviewer throughput: target >40 items/hour at confidence ≥ 0.85
- Rejection rate: target <8% overall (>8% indicates generation quality problem)
- First-30-day customer NPS: maintain >50 through human-review period

Flow 13 — Customer Health Dashboard (Operations / CSM)

Persona: Operations Lead, CSM **Goal:** Surface at-risk customers before they churn; surface expansion-ready customers before competitors do.

Health dashboard

Customer Health	47 customers tracked
Filter: ▾ All tiers ▾ All states	Sort: ▾ Health (asc)
Plano Electric Pros (Pro, day 67) Health: 23 • Score plateau 4 weeks • Digest not opened 3 weeks • Approval queue untouched 14 days Suggested action: outreach call this week [View account] [Trigger save flow] [Snooze]	
Reliable Electric (Lite, day 52) Health: 64 • Score lift solid (12 → 31) • Hit 25-prompt cap • 5 service-page interest events Suggested action: upgrade-to-Pro nudge [View] [Send upgrade prompt] [Schedule call]	
Lone Star Electric (Pro, day 41) Health: 89 • Score 12 → 47 (+35) • Autopilot enabled day 35 • Sent 2 referrals	

		Suggested action: case study request		
		[View] [Request case study]		

Health Score components

Health Score (0-100) =
0.30 × Score Trajectory (week-over-week Δ in AI Visibility Score)
+ 0.20 × Engagement (digest opens, queue interactions, login frequency)
+ 0.20 × Approval Quality (approval rate, edit rate, revert rate)
+ 0.15 × Expansion Signals (cap hits, gated feature views)
+ 0.10 × Support Load (negative tickets, escalations – inverse)
+ 0.05 × Tenure (longer = healthier, all else equal)

Triggered alerts

Trigger	Alert routing
Health drops below 40	CSM Slack channel + email
Health drops below 25	Ops lead immediate alert
Score plateau 4+ weeks	CSM action queue
Digest unread 3+ weeks	Auto-pause email frequency, send "are you still there?"
Cap-hit on Lite	Auto-trigger upgrade prompt + sales motion
90+ days, score >50, autopilot enabled	Case study request queue

Key UX decisions

- **Suggested action per row.** Internal team should not have to interpret data; dashboard prescribes the next move.
- **Color-coded urgency.** Same visual language as customer-facing app — single-color status.
- **Bulk actions.** Triage 50 customers at a time without 50 clicks.
- **Snooze, not dismiss.** Keep the customer in view; we don't lose them.

Events to instrument

```
health_score_recomputed      { customer_id, score, components }
health_alert_fired           { customer_id, severity }
health_action_taken          { customer_id, action, csm_id }
health_action_outcome        { customer_id, outcome: 'saved' | 'churned' |
                              'upgraded' }
```

Success metrics

- Time-to-intervention on red customers: target <48h
- Save rate on red intervention: target >50%
- Expansion conversion on yellow-with-upgrade-signal: target >30%

Flow 14 — Customer Account View (CSM Deep-Dive)

Persona: CSM, **Support Goal:** Single page that gives full context on one customer for any conversation.

Account view layout

Lone Star Electric[Pro] Day 67

Mike Reynolds, Owner • mike@lonestar.com • 972-555-0142

— HEALTH —

Health: 89 (green, ↗)

AI Visibility Score: 47 (started 12, +35 in 67 days)

Tier: Pro \$499/mo • 41 days to renewal

Autopilot: ON since day 35

— TIMELINE —

Day 67 Score 47 (+2 wk)

Day 65 Sent referral to Bright Spark Electrical

Day 35 Autopilot enabled

Day 18 Tier upgrade Lite → Pro

Day 12 Connected ServiceTitan

Day 1 First fix shipped

Day 0 Signed up via free audit

— ACTIVITY (LAST 30 DAYS) —

Fixes shipped: 47
Approval rate: 96% Edit rate: 4% Revert rate: 0%
Digest open rate: 100% (4/4)
Login frequency: 3x/wk
— INTEGRATIONS —
✓ Google Business Profile (lone-star-electric)
✓ WordPress (lonestarelectric.com)
✓ ServiceTitan (review-pump live)
× Slack (not connected)
— NOTES & FLAGS —
• Asked about multi-location during signup (1 location now, considering adding McKinney shop) – expansion target Q2
• Mentioned IEC chapter membership; potential case study
[Schedule call] [Send custom message] [Trigger upgrade]
[Export account data]

Key UX decisions

- **Single page.** No tabs, no drill-down. CSM grabs context in 30 seconds before a call.
- **Timeline view.** Shows trajectory, not just snapshot.
- **Notes inline.** CSMs can add free-text notes; persists with account.
- **Action buttons at the bottom.** Most-common actions one click away.

Flow 15 — Internal Onboarding Observation (First 30 Days)

Persona: CSM, Operations **Goal:** Watch every new customer's first 30 days like a hawk; intervene at any failure point.

Onboarding watch list

New Customers – First 30 Days	12 active
Day 1 (Today)	
✓ Lone Star Electric – first fix shipped 4h after signup	
▲ Plano Lighting Co – signed up 6h ago, no first fix yet	
[Send SMS reminder] [Owner is offline; queue for AM]	

Day 7

- ✓ Bright Spark – Score 8 → 19 (+11), digest opened
- ▲ Frisco Electric – Score flat, digest not opened
- [Schedule check-in call]

Day 14

- ✓ Reliable Electric – 7 fixes shipped, approval rate 100%
- Allen Electric Co – abandoned after onboarding
- [Send recovery sequence] [Trial ends in 3 days]

Day 30

- ✓ Wylie Watt – Score 14 → 38, ready for autopilot prompt
- ▲ McKinney Electric – Score lift modest; voice profile flag
- [Trigger voice retrain] [CSM call]

Daily ritual

- 9 AM standup: review overnight signups, day-7 / day-14 / day-30 milestones
- Anything in status gets intervention same-day
- Anything in ▲ status gets queued for AM call

Key UX decisions

- **Day-bucketed view.** First 30 days is a structured journey; surface where each customer is.
- **Proactive, not reactive.** CSM hits the SMS reminder *before* the customer churns, not after.
- **Auto-recovery sequences for abandoners.** Failed onboarding doesn't mean lost customer — it means triggered nurture.

Flow 16 — Support Inbox

Persona: Support, Tier-1 **Goal:** Resolve incoming tickets fast with shared macros and full account context.

Support inbox

Support Inbox

12 open • 3 priority

Lone Star – "Schema script broke our site!"
Reported 18 min ago | Owner: Mike R | Tier: Pro
[Open] [Assign to L2 (technical)]

Frisco Electric – "How do I add another service area?"
Reported 2h ago | Owner: Sarah T | Tier: Lite
Macro suggested: "Add service area" (94% match)
[Use macro] [Open]

Bright Spark – "Love the score lift"
Macro suggested: "Thank-you + case-study ask"

Macro library

- "Schema script install help" — links to walkthrough video
- "Add service area" — one-click flow that updates customer's service-area config
- "Score plateau explainer" — pre-written customer-facing explanation
- "Refund/cancel" — billing flow trigger
- "Thank-you + case-study ask" — for happy customers

Routing rules

- Billing/account/onboarding → Tier-1
- Schema/CMS/integration → Tier-2 (technical)
- Content quality/voice issues → Content Reviewer team
- Anything from a churn-risk account → CSM

Key UX decisions

- **Macro suggestion based on ticket content.** AI categorizes and suggests responses.
 - **Account context inline.** Customer's tier, score, day-in-product visible without clicking.
 - **Routing is automatic.** Tickets land where they should; manual reassignment is rare.
-

Flow 17 — Agency White-Label Admin (V2)

Persona: Marketing agency reselling Beacon **Goal:** Manage 10-100+ contractor sub-accounts under a single agency master.

Agency dashboard

ScorpionDigital – Agency Admin34 sub-accounts

Aggregate this week:

Avg Score lift: +6.2

Total fixes shipped: 312

Customers in red health: 2

Customers ready for QBR: 7

Sub-account	Score	Δ wk	Health	Action
Lone Star Electric	47	+2	89	–
Plano Electric Pros	31	0	23	▲ INTERVENE
Bright Spark Electrical	52	+4	77	–
Reliable Electric	41	+3	81	–
...				

[Bulk: run radar]

[Bulk: digest test]

[Add account]

Branding: ScorpionDigital logo on customer reports

Custom domain: insights.scorpiondigital.com ✓ Verified

Reseller pricing: \$399/sub-account flat (Pro features)

Key UX decisions

- **Aggregate view first, drill-down second.** Agency owners want signal across the portfolio.
- **Branding controls in the admin.** White-label is an actual feature, not a toggle.
- **Reseller pricing visible.** Agency knows its margin at all times.

PART III — CROSS-FLOW CONSIDERATIONS

Used everywhere a publish action is committed.

Approve and publish?

This will publish 6 FAQ pairs to:

- lonestarelectric.com/panel-up
- Google Business Profile Q&A

Goes live in ~5 minutes.

You can revert in one tap anytime.

[Cancel] [Approve & Publish]

Shared Pattern: Empty States

Empty states are designed as engagement opportunities, not error messages.

[Action Queue empty]

☐ All caught up

We're scanning AI engines now.

Next batch of fixes lands tomorrow morning by 7 AM.

[See last week's wins →]

Shared Pattern: Error & Fallback

All errors are user-grade, never developer-grade.

[GBP API timeout]


Hmm – Google's not responding.
We'll retry in a few minutes.

Your fix is queued and will publish
as soon as Google's back.

[Got it]

Cross-Flow PLG Architecture

How the customer-facing flows connect to drive PLG outcomes:



Instrumentation Plan Summary

Every event captured in this document maps to one of four buckets:

Bucket	Used for
Acquisition	Audit funnel, signup, source attribution
Activation	Onboarding completion, first-fix-shipped, time-to-aha
Engagement	Digest opens, queue interactions, login frequency
Conversion / Expansion	Upgrade clicks, referrals sent, autopilot enabled
Retention / Churn	Cancellation reasons, win-back funnel, save rates

A single event schema (`(event_name, customer_id, tier, state, properties{})`) is used everywhere. Stored in a warehouse (BigQuery or Snowflake) and surfaced via a single dashboard for the team.

Build Order for V0.1 (Aligned with PRD MVP Weeks 1-8)

Week	Customer-facing flows shipped	Internal flows shipped
1	— (skeleton + auth)	—
2	Audit (Flow 1, basic)	—
3	Audit polished + email capture	—
4	Onboarding wizard (Flow 2)	Content Review Queue (Flow 12) — minimal
5	Action Queue + Approval (Flow 5)	Onboarding Watch (Flow 15)
6	First-Fix Shipping + GBP integration	Customer Account View (Flow 14)
7	Weekly Digest (Flow 4)	Health Dashboard (Flow 13) — basic
8	Polish + bug bash	Support Inbox (Flow 16) — minimal

Deferred to V1 (weeks 9-20):

- Autopilot graduation (Flow 6)
- Score plateau intervention (Flow 7)
- Team invite (Flow 8)
- Tier upgrade modal (Flow 9)
- Referral system (Flow 10)
- Cancellation polish (Flow 11)

Deferred to V2 (months 6+):

- Agency white-label (Flow 17)
- Score certificate
- Multi-location consolidated view

— *End of v0.1. Next revision after first 10 design partners complete onboarding (week 6 of MVP).*